

# Matteo Bunino

## Digital Twin & Computing Engineer

Email: [matteo.bunino@cern.ch](mailto:matteo.bunino@cern.ch)

Web portfolio: [matbun.github.io](https://matbun.github.io)

### QUALIFICATIONS

**M.Sc. Data Science & Engineering** - Politecnico di Torino & EURECOM (double degree), 2022 – Grade: 110 cum laude / 110

**B.Sc. Computer Engineering** - Politecnico di Torino, 2019 – Grade: 108 / 110

### KEY STRENGTHS

Digital twin design and development for scientific applications • Requirements capture, system architecture and design (HPC/AI and digital twins) • Deployment of AI-driven digital twins on heterogeneous HPC • Distributed AI workflows (PyTorch, TensorFlow, Horovod, DeepSpeed, Ray) • Energy benchmarking on GPUs (Nvidia, AMD) and Arm • Hyperparameter optimization • HPC and cloud workload orchestration (SLURM + Singularity, Kubernetes + Docker) • Cloud–HPC observability and energy optimization (Prometheus, Kepler, PUE analysis) • Software development and coordination in EU projects ([interTwin Task 6.5](#) lead, ODISSEE contributor) • Technical writing (e.g., EU deliverables), conference presentations, event organization.

### TECHNICAL SKILLS

**Digital twin frameworks:** Nvidia [Omniverse](#) for 3D visualizations, [ExaDigiT](#) (ORNL)

**HPC & cloud:** [itwinai](#), SLURM, [interLink](#), Kubernetes, Prometheus, Kepler, Grafana

**AI and MLOps:** PyTorch, TensorFlow, Horovod, DeepSpeed, Ray, Kubeflow, MLflow

**Containers & CI/CD:** Docker, Singularity/Aptainer, Dagger, GitHub Actions

**Programming & Scripting:** Python, Go, C++, Bash

### WORK EXPERIENCE

**Fellow** (Jan 2023 - present) - *CERN, Geneva, Switzerland*

- Main developer of [itwinai](#) in [interTwin](#) (Task 6.5 lead), an abstraction layer for distributed AI on HPC systems, combining hyperparameter optimization and reducing engineering overheads for scientists; validated itwinai on HPC by integrating four scientific digital twin use cases from physics (e.g., MLPF) and three from Earth observation; optimized use case code scalability on HPC.
- CERN openlab digital twin expert: led technical discussions with Nvidia, Google, and Imperial College; coordinated two digital twin projects: digital twin of the LHC infrastructure with Nvidia Omniverse, and digital twin for data quality monitoring at CMS ECAL in collaboration with Imperial College.

- Digital twin HPC benchmarking: measured performance and energy consumption of scientific digital twin AI workloads across different European HPC centres: LUMI (AMD MI250X), Deucalion (Arm Fujitsu A64FX), Jülich (Nvidia A100/V100), and Vega (Nvidia A100).
- Cloud-HPC integration via [interLink](#), enabling large-scale AI workloads using the Ray Kubernetes operator and cross-infrastructure container CI/CD with Dagger (presented at [KubeCon EU 2025](#)).
- Contributed to EC-funded [ODISSEE](#) project by collecting use case requirements and assessing digital twin frameworks for modelling and optimizing data centre operations (e.g., [ExaDigiT](#)); assessed Energy Aware Runtime ([EAR](#)) for energy-aware job execution; liaised with project partners (SURF, LHCb, NextSilicon, SiPearl).
- Visualization of energy and resource utilization telemetry (Prometheus) on Grafana dashboards.
- 3D visualizations with Nvidia Omniverse for interactive digital twins.
- Digital twin events: organized a Birds-of-a-Feather (BoF) session on digital twins at ISC 2025 (speakers: Nvidia, SURF, CINES, ECMWF, CERN); presented work on cloud-HPC integration and itwinai at ISC 2024, CHEP 2024, and PASC 2025.
- Supervised 4 summer students and 3 technical students in digital twin projects; supported the recruiting of new fellows.
- EC-funded projects: [interTwin](#) (2023 - 2025), [ODISSEE](#) (2025 - present)

### **Cyber Security Data Scientist** (Jun 2022 - Dec 2022) - *Huawei, Munich, Germany*

Co-designed and co-developed a malware analysis framework based on large language models (LLMs); Scaled AI workloads to Huawei Cloud. Filed periodic technical progress reports; Graph machine learning; Natural language processing.

### **Research Intern** (Sep 2021 - Feb 2022) - *Huawei, Munich, Germany*

Developed a proof-of-concept that uses Reinforcement Learning (RL) to automate the reverse engineering of evasive malware, improving the efficiency and accuracy of dynamic malware analysis. Find my thesis [here](#).

## **SELECTED TALKS AND PUBLICATIONS**

- Bunino M. et al. “itwinai: Enabling Scalable AI Workflows on HPC for Digital Twins in Science” - EPJC, 2025, DOI: <https://doi.org/10.1051/epjconf/202533701361>
- Bunino M. et al. “[Testing AI Containers for Digital Twins in Science: a Cloud-HPC Workflow](#)” - KubeCon EU 2025
- BoF organiser & speaker, [Shaping the Future of Scientific Digital Twins: Tales from Physics, Climate Research, and Data Centre Operations](#), ISC 2025
- BoF speaker, [Synergistically Integrating HPC and Cloud](#), ISC 2024

## **LANGUAGES**

**Italian:** Native

**English:** Fluent

**French:** Conversational